

---

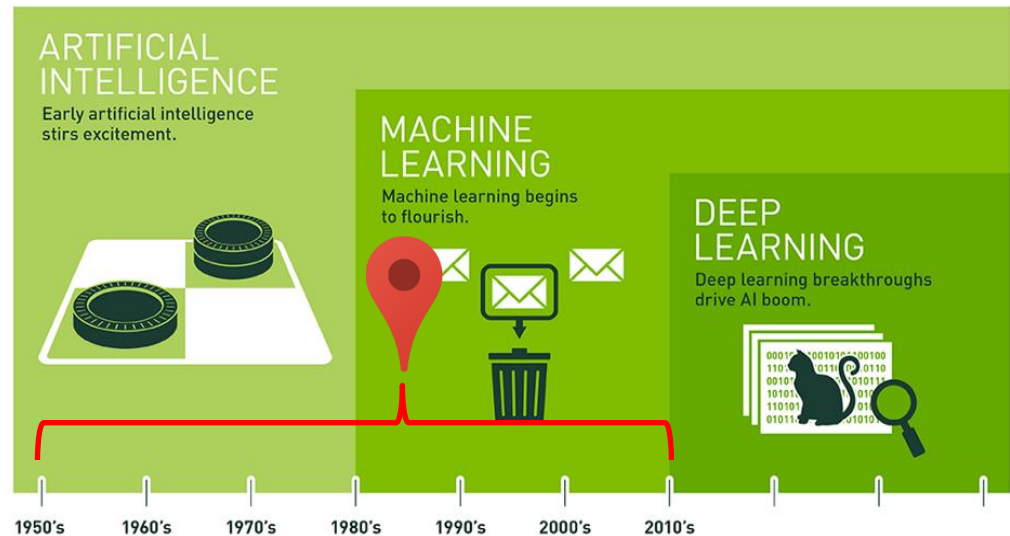
# Artificial Intelligence: Introduction to Neural Networks

Perceptron, Backpropagation

# Today

## ■ Neural Networks

- Perceptrons
- Backpropagation



# Neural Networks


- Radically different approach to reasoning and learning
- Inspired by biology
  - the neurons in the human brain
- Set of many simple processing units (neurons) connected together
- Behavior of each neuron is very simple
  - but a collection of neurons can have sophisticated behavior and can be used for complex tasks
- In a neural network, the behavior depends on weights on the connection between the neurons
- The weights will be learned given training data

# Biological Neurons

- Human brain =
  - 100 billion neurons
  - each neuron may be connected to 10,000 other neurons
  - passing signals to each other via 1,000 trillion **synapses**



- A neuron is made of:
  - **Dendrites**: filaments that provide input to the neuron
  - **Axon**: sends an output signal
  - **Synapses**: connection with other neurons - releases neurotransmitters to other neurons

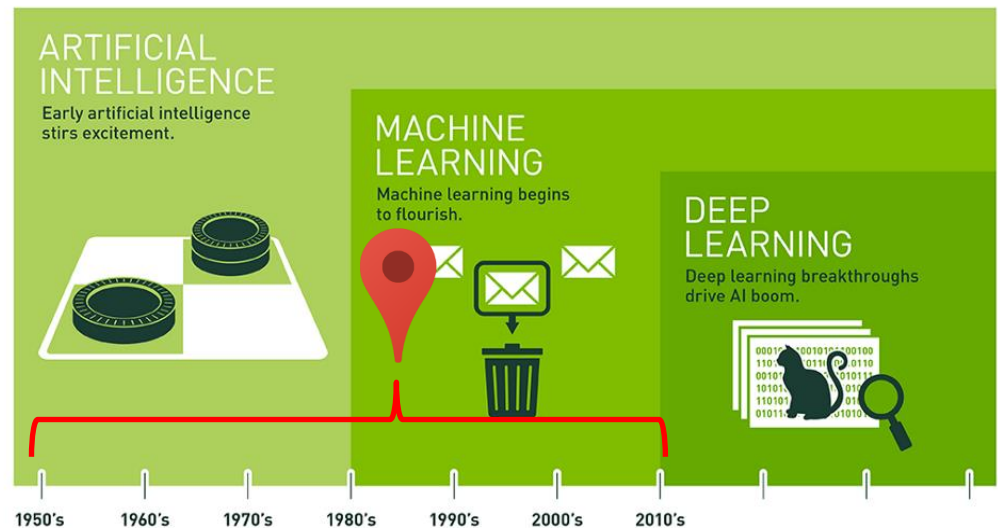


# Behavior of a Neuron

- A neuron receives inputs from its neighbors
- If enough inputs are received at the same time:
  - the neuron is **activated**
  - and **fires** an output to its neighbors
- Repeated firings across a synapse increases its sensitivity and the future likelihood of its firing
- If a particular stimulus repeatedly causes activity in a group of neurons, they become strongly associated

# Today

- Neural Network
  - Perceptrons
  - Backpropagation



# Feature Vector Representation

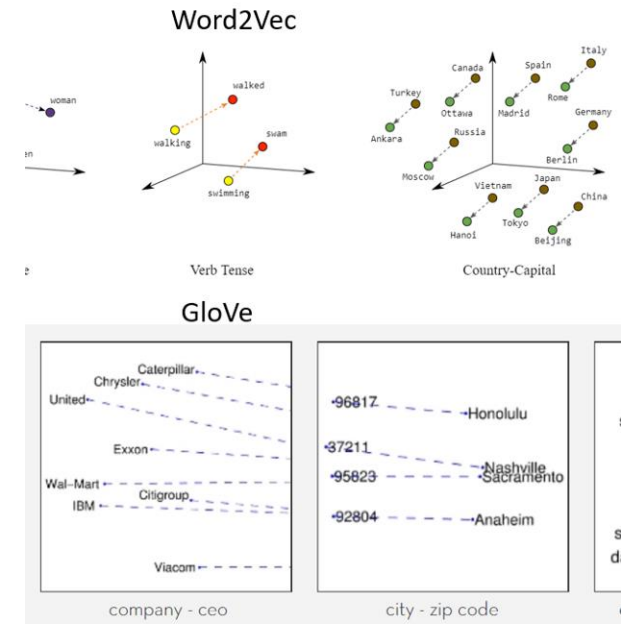
- Sources of Feature Vector  $x$ 
  - ❑ Encoded image
  - ❑ Tabulated data
  - ❑ Embedded words
  - ❑ ...



Tabular Data

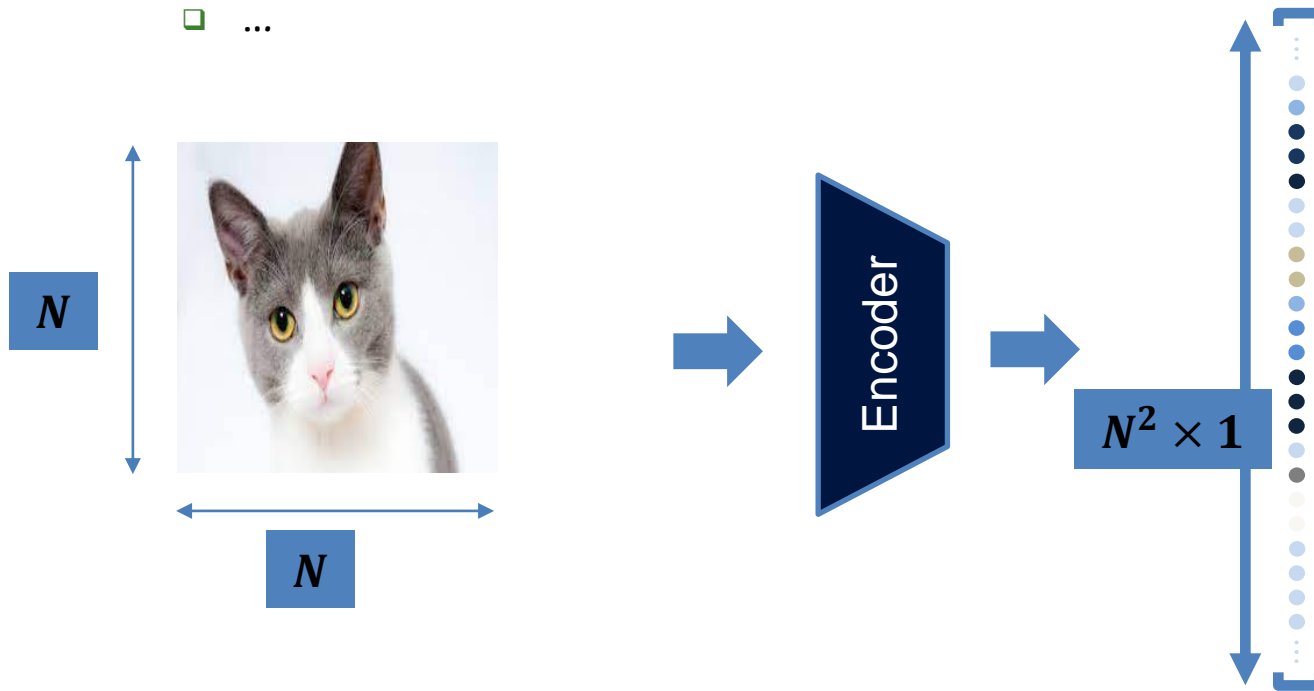
columns = attributes for those observations

Player	Minutes	Points	Rebounds	Assists
A	41	20	6	5
B	30	29	7	6
C	22	7	7	2
D	26	3	3	9
E	20	19	8	0
F	9	6	14	14
G	14	22	8	3
I	22	36	0	9
J	34	8	1	3



# Feature Vector Representation


- Sources of Feature Vector  $x$ 
  - Encoded image
  - Tabulated data
  - Embedded words
  - ...





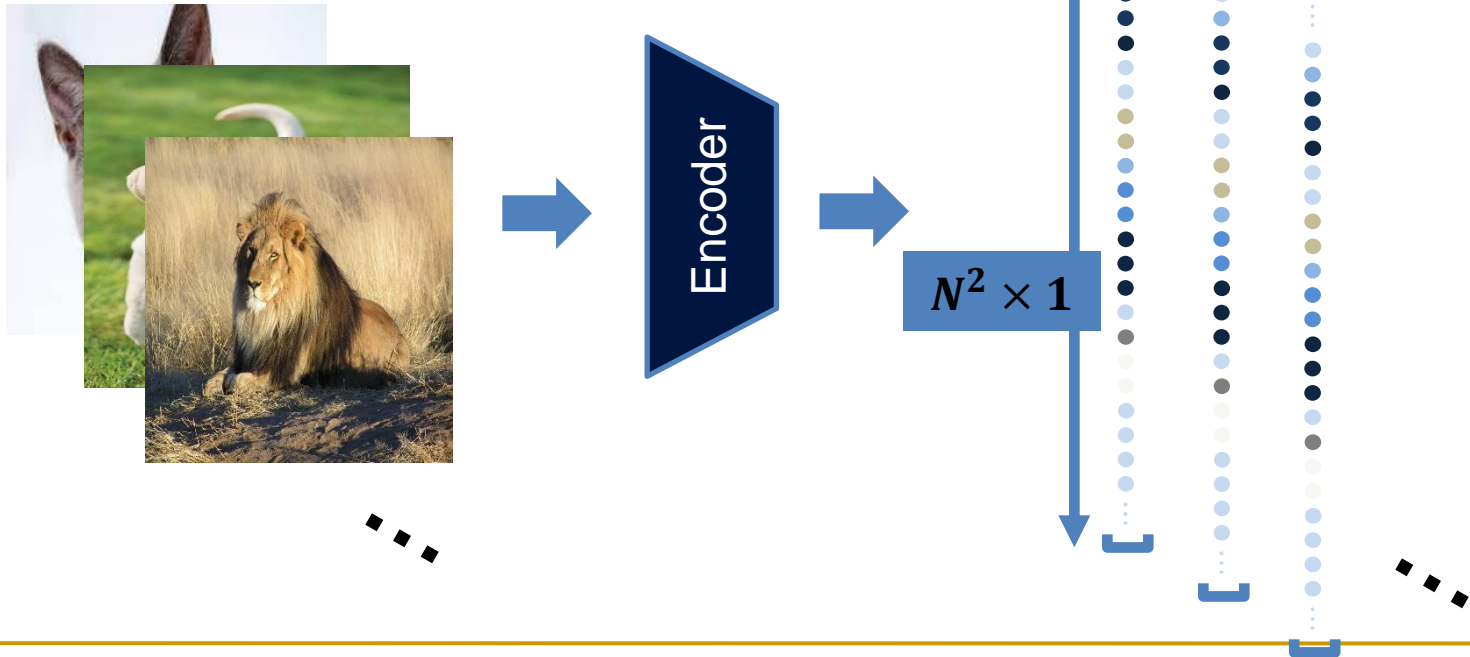
# Feature Vector Representation

- Sources of Feature Vector  $x$ 
  - Encoded image
  - Tabulated data
  - Embedded words
  - ...



# Feature Vector Representation


- Sources of Feature Vector  $x$ 
  - Encoded image
  - Tabulated data
  - Embedded words
  - ...



# A Perceptron Network

- **Goal:** Map *Input Feature Vector*  $\mathbf{x}$  into *Output Feature Vector*  $\mathbf{y}$


$$\sum_{i=1}^N w_i x_i$$



# A Perceptron Network

- **Goal:** Map *Input Feature Vector*  $\mathbf{x}$  into *Output Feature Vector*  $\mathbf{y}$


$$\sum_{i=1}^N w_i x_i + b_j$$



# A Perceptron Network

- **Goal:** Map *Input Feature Vector*  $\mathbf{x}$  into *Output Feature Vector*  $\mathbf{y}$


$$y_j = f \left( \sum_{i=1}^N w_i x_i + b_j \right)$$



# A Perceptron Network

- **Goal:** Map *Input Feature Vector*  $\mathbf{x}$  into *Output Feature Vector*  $\mathbf{y}$

$$y_j = f \left( \sum_{i=1}^N w_i x_i + b_j \right) = f (w^T x + b_j)$$



# A Perceptron Network


- **Goal:** Map Input Feature Vector  $\mathbf{x}$  into Output Feature Vector  $\mathbf{y}$

$$y_j = f \left( \sum_{i=1}^N w_i x_i + b_j \right) = f$$

Activation functions


**Sigmoid**

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




**tanh**

$$\tanh(x)$$




**ReLU**

$$\max(0, x)$$



**Leaky ReLU**

$$\max(0.1x, x)$$

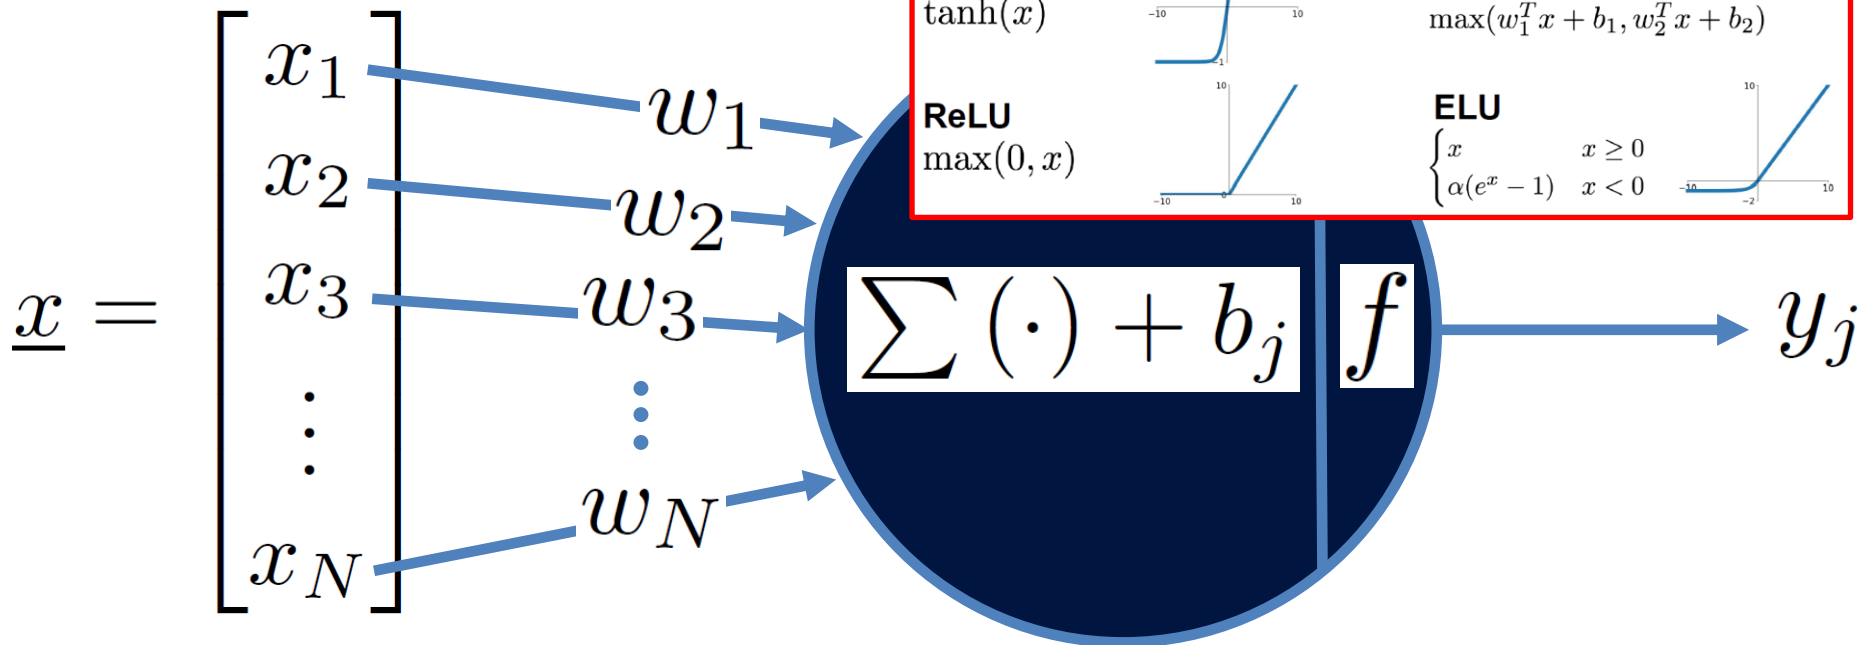



**Maxout**

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

**ELU**


$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$




# A Perceptron Network

- **Goal:** Map Input

$$y_j = f \left( \sum_{i=1}^N$$




ure Vector  $y$







# A Perceptron Network




# A Perceptron Network




# A Perceptron Network




# A Perceptron Network




# A Perceptron Network




# A Perceptron Network



# A Perceptron Network



# A Perceptron Network



$$\underline{y} = f(W^T \underline{x} + \underline{b})$$

$\underline{x}$ : input feature vector  $N \times 1$

$\underline{y}$ : input feature vector  $M \times 1$


$\underline{W}$ : weight matrix  $M \times N$

$\underline{b}$ : bias vector  $M \times 1$

$f(\cdot)$ : activation function




# A Perceptron Network



$$\leftarrow \underline{x}_1 = f(W_0^T \underline{x} + \underline{b}_0)$$

# A Perceptron Network




←  $\underline{x}_1 = f(W_0^T \underline{x} + \underline{b}_0)$

←  $\underline{y} = f(W_1^T \underline{x}_1 + \underline{b}_1)$




# Fully Connected (FC) Network



# Neural Network: FC Layer Design


- First layer design: map input feature vector into smaller vector



# Neural Network: FC Layer Design

- First layer design: activate features using non-linear activation

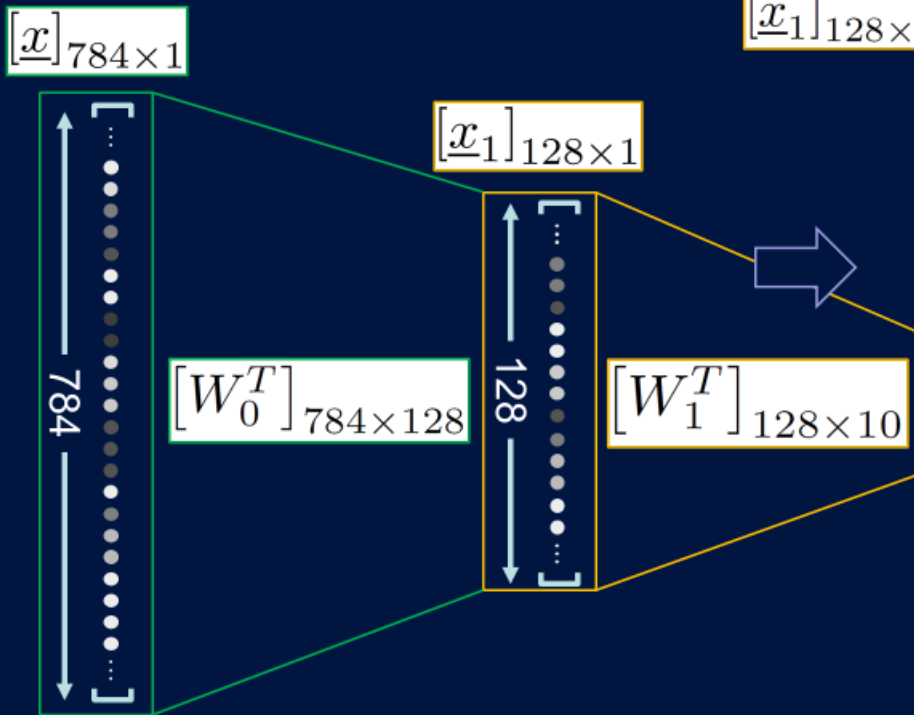
$$[\underline{x}_1]_{128 \times 1} = f(W_0^T \underline{x} + \underline{b}_0)$$



## Neural Network: FC Layer Design


- Second layer design: map input feature vector into #Classes

$$[\underline{x}_1]_{128 \times 1} = f(W_0^T \underline{x} + \underline{b}_0)$$



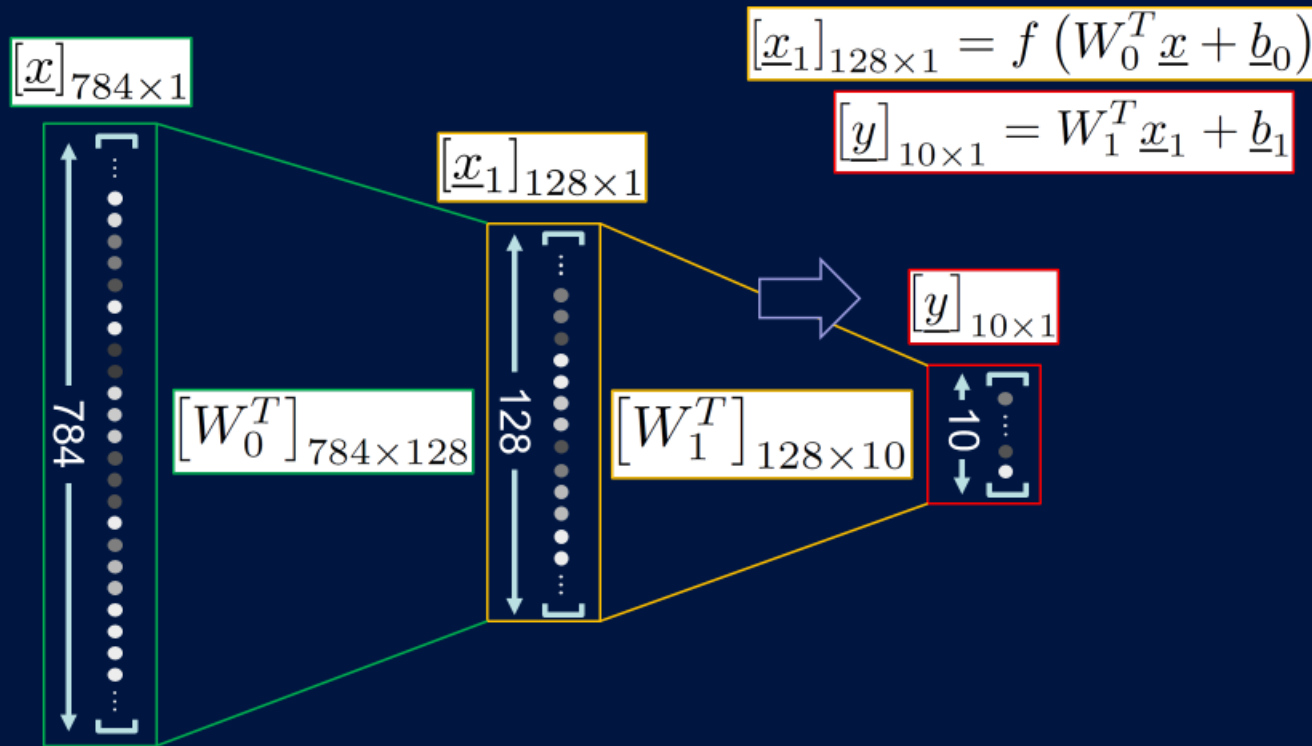
# Neural Network: FC Layer Design

- Second layer design: no activation is required



# Neural Network: FC Layer Design


**Quiz:** What is the #Learnable Parameters?






# Neural Network: FC Layer Design

#Learnable Parameters =  $784 \times 128 + 128 + 128 \times 10 + 10 = 101,770$



# Neural Network: FC Layer Design

Apply Softmax regression model to map output classes in range between [0,1]



# Fully Connected NN: Error-Gradient Backpropagation

- To train NN, the gradient of error-loss is calculated with respect to each learnable parameter in the network

$$\frac{\partial \epsilon}{\partial w_{0,i,j}}, \frac{\partial \epsilon}{\partial b_{0,j}}, \frac{\partial \epsilon}{\partial w_{1,i,j}}, \frac{\partial \epsilon}{\partial b_{1,j}}$$

# Fully Connected NN: Error-Gradient Backpropagation

- To train NN, the gradient of error-loss is calculated with respect to each learnable parameter in the network

$$\frac{\partial \epsilon}{\partial w_{0,i,j}}, \frac{\partial \epsilon}{\partial b_{0,j}}, \frac{\partial \epsilon}{\partial w_{1,i,j}}, \frac{\partial \epsilon}{\partial b_{1,j}}$$

Rumelhart et al., 1986  
introduced Backpropagation  
for training NN

- Error gradients are used to update the network parameters in iterative minimization using gradient descent

$$\text{e.g. } w_{0,i,j}^{k+1} \leftarrow w_{0,i,j}^k - \eta \frac{\partial \epsilon}{\partial w_{0,i,j}^{k+1}}$$



The Edward S. Rogers Sr. Department  
of Electrical & Computer Engineering  
UNIVERSITY OF TORONTO

FACULTY  
OF APPLIED  
SCIENCE &  
ENGINEERING


# Fully Connected NN: Error-Gradient Backpropagation

- To train NN, the gradient of error-loss is calculated with respect to each learnable parameter in the network

$$\frac{\partial \epsilon}{\partial w_{0,i,j}}, \frac{\partial \epsilon}{\partial b_{0,j}}, \frac{\partial \epsilon}{\partial w_{1,i,j}}, \frac{\partial \epsilon}{\partial b_{1,j}}$$

- Error gradients are used to update the network parameters in iterative minimization using gradient descent

$$\text{e.g. } w_{0,i,j}^{k+1} \leftarrow w_{0,i,j}^k - \eta \frac{\partial \epsilon}{\partial w_{0,i,j}^{k+1}}$$



The Edward S. Rogers Sr. Department  
of Electrical & Computer Engineering  
UNIVERSITY OF TORONTO

FACULTY  
OF APPLIED  
SCIENCE &  
ENGINEERING

# Fully Connected NN: Error-Gradient Backpropagation

- To train NN, the gradient of error-loss is calculated with respect to each learnable parameter in the network

$$\frac{\partial \epsilon}{\partial w_{0,i,j}}, \frac{\partial \epsilon}{\partial b_0}$$


- Error gradient iterative mini

$$\text{e.g. } w_{0,i,j}^{k+1} \leftarrow$$

```
90 iteration_p = 0
91 for p in group['params']:
92     if p.grad is None:
93         iteration_p += 1
94         continue
95     d_p = p.grad.data
96     if weight_decay != 0:
97         d_p.add_(weight_decay, p.data)
98     if momentum != 0:
99         param_state = self.state[p]
100         if 'momentum_buffer' not in param_state:
101             buf = param_state['momentum_buffer'] = torch.clone(d_p).detach()
102         else:
103             buf = param_state['momentum_buffer']
104             buf.mul_(momentum).add_(1 - dampening, d_p)
105         if nesterov:
106             d_p = d_p.add(momentum, buf)
107         else:
108             d_p = buf
109
110     p.data.add_(-group['lr'], d_p)
111     iteration_p += 1
112
113 return loss
```

## Fully Connected NN: Calculating Error-Gradients using Chain-Rule

- Classification sub-module



$$\epsilon(\hat{y}, t) = - \sum_{k=1}^{10} t_k \ln \hat{y}_k = -\underline{t}^T \ln \underline{\hat{y}}$$




The Edward S. Rogers Sr. Department  
of Electrical & Computer Engineering  
UNIVERSITY OF TORONTO

FACULTY  
OF APPLIED  
SCIENCE &  
ENGINEERING

## Fully Connected NN: Calculating Error-Gradients using Chain-Rule

- Classification sub-module



$$\epsilon(\hat{y}, t) = - \sum_{k=1}^{10} t_k \ln \hat{y}_k = - \underline{t}^T \ln \underline{\hat{y}}$$

$$\frac{\partial \epsilon}{\partial \hat{y}_j} = - \frac{t_j}{\hat{y}_j}$$



# Fully Connected NN: Calculating Error-Gradients using Chain-Rule

- Probability sub-module



$$\hat{y}_j = \frac{\exp y_j}{\sum_{k=1}^{10} \exp y_k}$$




The Edward S. Rogers Sr. Department  
of Electrical & Computer Engineering  
UNIVERSITY OF TORONTO

FACULTY  
OF APPLIED  
SCIENCE &  
ENGINEERING

## Fully Connected NN: Calculating Error-Gradients using Chain-Rule

- Probability sub-module




$$\hat{y}_j = \frac{\exp y_j}{\sum_{k=1}^{10} \exp y_k}$$

$$\frac{\partial \epsilon}{\partial y_j} = \frac{\partial \epsilon}{\partial \hat{y}_j} \cdot \frac{\partial \hat{y}_j}{\partial y_j} = \frac{\partial \epsilon}{\partial \hat{y}_j} \cdot \hat{y}_j \cdot (1 - \hat{y}_j)$$

# Fully Connected NN: Calculating Error-Gradients using Chain-Rule

- FC Layer sub-module

$$y_j = \sum_{k=1}^{128} W_{1,k,j} x_k + b_{1,j}$$




# Fully Connected NN: Calculating Error-Gradients using Chain-Rule

- FC Layer sub-module

$$y_j = \sum_{k=1}^{128} W_{1,k,j} x_k + b_{1,j}$$

$$\frac{\partial \epsilon}{\partial b_{1,i}} = \frac{\partial \epsilon}{\partial y_j} \cdot \frac{\partial y_j}{\partial b_{1,i}} = \frac{\partial \epsilon}{\partial y_j} \cdot 1$$



The Edward S. Rogers Sr. Department  
of Electrical & Computer Engineering  
UNIVERSITY OF TORONTO

FACULTY  
OF APPLIED  
SCIENCE &  
ENGINEERING


# Fully Connected NN: Calculating Error-Gradients using Chain-Rule

- FC Layer sub-module

$$y_j = \sum_{k=1}^{128} W_{1,k,j} x_k + b_{1,j}$$

$$\frac{\partial \epsilon}{\partial b_{1,i}} = \frac{\partial \epsilon}{\partial y_j} \cdot \frac{\partial y_j}{\partial b_{1,i}} = \frac{\partial \epsilon}{\partial y_j} \cdot 1$$

$$\frac{\partial \epsilon}{\partial W_{1,i,j}} = \frac{\partial \epsilon}{\partial y_j} \cdot \frac{\partial y_j}{\partial W_{1,i,j}} = \frac{\partial \epsilon}{\partial y_j} \cdot x_i$$



# Fully Connected NN: Calculating Error-Gradients using Chain-Rule


- FC Layer sub-module

$$y_j = \sum_{k=1}^{128} W_{1,k,j} x_k + b_{1,j}$$

$$\frac{\partial \epsilon}{\partial b_{1,i}} = \frac{\partial \epsilon}{\partial y_j} \cdot \frac{\partial y_j}{\partial b_{1,i}} = \frac{\partial \epsilon}{\partial y_j} \cdot 1$$

$$\frac{\partial \epsilon}{\partial W_{1,i,j}} = \frac{\partial \epsilon}{\partial y_j} \cdot \frac{\partial y_j}{\partial W_{1,i,j}} = \frac{\partial \epsilon}{\partial y_j} \cdot x_i$$

$$\frac{\partial \epsilon}{\partial x_i} = \frac{\partial \epsilon}{\partial y_j} \cdot \frac{\partial y_j}{\partial x_i} = \frac{\partial \epsilon}{\partial y_j} \cdot W_{1,i,j}$$




# Applications of Neural Networks

- Handwritten digit recognition
  - Training set = set of handwritten digits (0...9)
  - Task: given a bitmap, determine what digit it represents
  - Input: 1 feature for each pixel of the bitmap
  - Output: 1 output unit for each possible character (only 1 should be activated)
  - After training, network should work for fonts (handwriting) never encountered

- Related pattern recognition applications:

- recognize postal codes
- recognize signatures
- ...




# Applications of Neural Networks

- Speech synthesis

- Learning to pronounce English words
- Difficult task for a rule-based system because English pronunciation is highly irregular
- Examples:
  - letter "c" can be pronounced [k] (*cat*) or [s] (*cents*)
  - *Woman* vs *Women*
- NETtalk:
  - uses the context and the letters around a letter to learn how to pronounce a letter
  - Input: letter and its surrounding letters
  - Output: phoneme



# NETtalk Architecture



Ex: *a cat* → *c is pronounced K*

- Network is made of 3 layers of units
- input unit corresponds to a 7 character window in the text
- each position in the window is represented by 29 input units (26 letters + 3 for punctuation and spaces)
- 26 output units - one for each possible phoneme

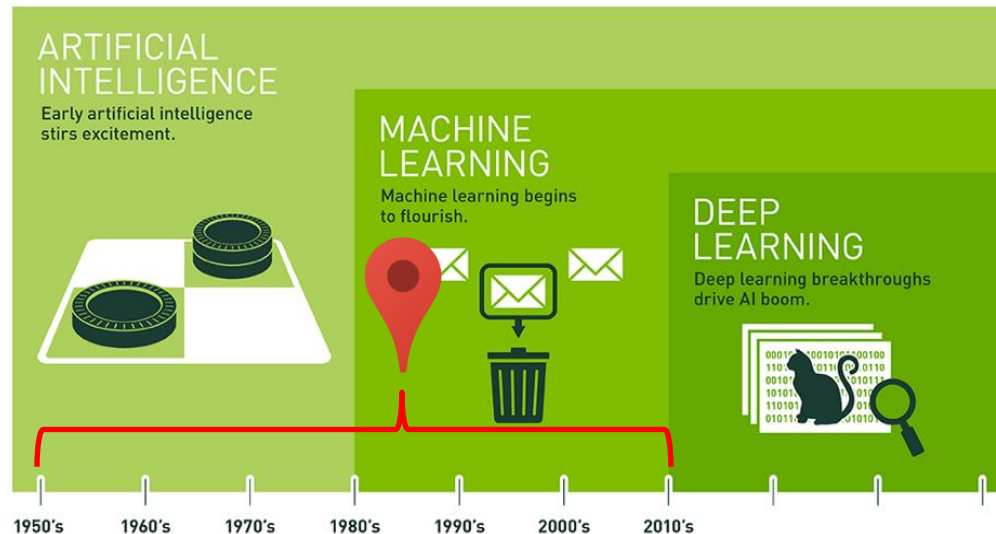
Listen to the output through iterations: <https://www.youtube.com/watch?v=gakJlr3GecE>

# Neural Networks

- Disadvantage:
  - result is not easy to understand by humans (set of weights compared to decision tree)... it is a black box
- Advantage:
  - robust to noise in the input (small changes in input do not normally cause a change in output) and graceful degradation

# Today

- Introduction to Neural Networks
  - Perceptrons
  - Backpropagation



YOU ARE HERE!

THE END!